

StockTracker Implementation Report

Time Complexity Requirements

The project specifies the following operations with their required time complexities:

- `insert_new_stock(x, p)`: Insert a stock with ID x and price p , $O(\log n)$.
- `update_price(x, p)`: Update the price of stock x to p , $O(\log n)$.
- `increase_volume(x, vinc)`: Increase the trading volume of stock x by $vinc$, $O(\log n)$.
- `lookup_by_id(x)`: Return the price and volume of stock x , $O(1)$.
- `price_range(p1, p2)`: Return IDs of stocks with prices in $[p1, p2]$, $O(\log n + k)$, where k is the number of stocks in the range.
- `max_vol()`: Return the maximum trading volume and its stock ID, $O(1)$ for retrieval, $O(\log n)$ for heap cleanup.

Here, n is the number of stocks stored.

Data Structures and Algorithms

To meet these time guarantees, the StockTracker class employs a combination of data structures: a hash table, a Binary Search Tree (BST), a defaultdict for price-to-ID mapping, and a max heap with UIDs for volume tracking. Below, each operation is analyzed with its supporting data structure and algorithm.

1. Hash Table (`self.stocks`)

- **Purpose:** Stores stock information (price and volume) for each stock ID.
- **Implementation:** A Python dictionary (`self.stocks`), where keys are stock IDs (x) and values are lists [price, volume].
- **Operations:**
 - **Lookup:** Checking if a stock exists (x in `self.stocks`) or retrieving its data (`self.stocks[x]`) is $O(1)$ on average due to the hash table's efficient key-value mapping.

- **Insert/Update:** Adding or updating a stock (`self.stocks[x] = [p, 0]`) is O(1).
- **Usage:** Critical for `lookup_by_id`, `insert_new_stock`, `update_price`, and `increase_volume` to achieve fast access and updates.

2. Binary Search Tree (`self.price_sorted`)

- **Purpose:** Maintains stock IDs sorted by price for efficient range queries.
- **Implementation:** A custom `BinarySearchTree` class with `BSTNode` objects, each storing a price and id. The BST supports:
 - **Insert:** Adds a node with price and id, O(log n) in a balanced BST (average case).
 - **Delete:** Removes a node with price and id, O(log n).
 - **Range Query:** Retrieves IDs with prices in $[p_1, p_2]$ using in-order traversal, O(log n + k).
- **Algorithm Details:**
 - **Insert:** Traverses the tree based on price (and id for ties), inserting a new node at the appropriate leaf.
 - **Delete:** Finds the node, handles three cases (no children, one child, two children), and uses the minimum node in the right subtree as a successor for two-child cases.
 - **Range Query:** Performs an in-order traversal, pruning branches outside $[p_1, p_2]$, collecting IDs within the range.
- **Usage:** Supports `insert_new_stock`, `update_price` (delete and insert), and `price_range`.

3. Defaultdict (`self.price_to_ids`)

- **Purpose:** Maps prices to sets of stock IDs to handle multiple stocks with the same price efficiently.
- **Implementation:** A `defaultdict(set)` where keys are prices and values are sets of stock IDs.
- **Operations:**
 - **Add/Remove:** Adding an ID (`self.price_to_ids[p].add(x)`) or removing an ID (`self.price_to_ids[p].discard(x)`) is O(1).
 - **Cleanup:** Deleting an empty set (`del self.price_to_ids[old_price]`) is O(1).
- **Usage:** Used in `insert_new_stock` and `update_price` to track IDs per price, ensuring

efficient BST updates.

4. Max Heap (`self.volume_heap`) with UUIDs (`self.id_to_uuid`)

- **Purpose:** Tracks the stock with the maximum trading volume.
- **Implementation:** A Python list (`self.volume_heap`) used as a max heap via `heapq`, storing tuples (-volume, x, uuid). `self.id_to_uuid` maps stock IDs to their latest UUID to invalidate outdated heap entries.
- **Operations:**
 - **Push:** Adding a new volume entry (`heapq.heappush`) is $O(\log n)$.
 - **Peek/Cleanup:** Checking the top entry and removing invalid entries (due to volume updates) is $O(1)$ for peeking and $O(\log n)$ for popping invalid entries.
- **Algorithm Details:**
 - Each volume update generates a new UUID, pushing a new (-volume, x, uuid) tuple.
 - `max_vol` checks the heap's top entry, popping entries with outdated UUIDs or deleted stocks until a valid entry is found.
- **Usage:** Supports `increase_volume` (push) and `max_vol` (peek/cleanup).

Implementation in Source Code

The `stock_tracker.py` file implements the above data structures and algorithms as follows:

1. `insert_new_stock(self, x, p)` - $O(\log n)$

- **Code:**

```
def insert_new_stock(self, x, p):  
    if x in self.stocks:  
        return  
  
    self.stocks[x] = [p, 0]  
    self.price_sorted.insert(p, x)  
    self.price_to_ids[p].add(x)  
    unique_id = str(uuid.uuid4())  
    self.id_to_uuid[x] = unique_id  
    heapq.heappush(self.volume_heap, (0, x, unique_id))
```

- **Steps:**
 - Checks if x exists in self.stocks ($O(1)$).
 - Adds [p, 0] to self.stocks ($O(1)$).
 - Inserts (p, x) into the BST (self.price_sorted.insert, $O(\log n)$).
 - Adds x to self.price_to_ids[p] ($O(1)$).
 - Generates a UUID and pushes (0, x, uuid) to the max heap ($O(\log n)$).
- **Time Guarantee:** Dominated by BST insert and heap push, both $O(\log n)$.

2. update_price(self, x, p) - $O(\log n)$

- **Code:**

```
def update_price(self, x, p):
    if x not in self.stocks:
        return
    old_price, volume = self.stocks[x]
    self.stocks[x][0] = p
    self.price_to_ids[old_price].discard(x)
    if not self.price_to_ids[old_price]:
        del self.price_to_ids[old_price]
    self.price_sorted.delete(old_price, x)
    self.price_sorted.insert(p, x)
    self.price_to_ids[p].add(x)
```
- **Steps:**
 - Checks if x exists ($O(1)$).
 - Updates price in self.stocks ($O(1)$).
 - Removes x from self.price_to_ids[old_price] and cleans up if empty ($O(1)$).
 - Deletes (old_price, x) from BST ($O(\log n)$).
 - Inserts (p, x) into BST ($O(\log n)$).
 - Adds x to self.price_to_ids[p] ($O(1)$).
- **Time Guarantee:** Dominated by BST delete and insert, $O(\log n)$.

3. increase_volume(self, x, vinc) - $O(\log n)$

- **Code:**

```
def increase_volume(self, x, vinc):
    if x not in self.stocks:
```

```

        return

    self.stocks[x][1] += vinc

    new_volume = self.stocks[x][1]
    unique_id = str(uuid.uuid4())
    self.id_to_uuid[x] = unique_id
    heapq.heappush(self.volume_heap, (-new_volume, x, unique_id))

```

- **Steps:**
 - Checks if x exists ($O(1)$).
 - Increments volume in self.stocks ($O(1)$).
 - Generates a new UUID and pushes $(-\text{new_volume}, \text{x}, \text{uuid})$ to the max heap ($O(\log n)$).
- **Time Guarantee:** Dominated by heap push, $O(\log n)$.

4. `lookup_by_id(self, x)` - $O(1)$

- **Code:**

```
def lookup_by_id(self, x):
    if x not in self.stocks:
        return None
    return self.stocks[x]
```
- **Steps:**
 - Checks if x exists in self.stocks ($O(1)$).
 - Returns self.stocks[x] or None ($O(1)$).
- **Time Guarantee:** Hash table operations are $O(1)$ on average.

5. `price_range(self, p1, p2)` - $O(\log n + k)$

- **Code:**

```
def price_range(self, p1, p2):
    return self.price_sorted.range_query(p1, p2)
```
- **Steps:**
 - Calls `BinarySearchTree.range_query`, which:
 - Traverses the BST to the first node where $\text{price} \geq p1$ ($O(\log n)$).
 - Performs in-order traversal, collecting IDs where $p1 \leq \text{price} \leq p2$

(O(k)).

- Prunes branches where price > p2 (reduces unnecessary traversal).
- **Time Guarantee:** O(log n) to reach the range, O(k) to collect k IDs.

6. max_vol(self) - O(1) retrieval, O(log n) cleanup

- **Code:**

```
def max_vol(self):
```

```
    while self.volume_heap:  
        neg_vol, x, uuid = self.volume_heap[0]  
        if x not in self.stocks or self.id_to_uuid.get(x) != uuid:  
            heapq.heappop(self.volume_heap)  
            continue  
        return -neg_vol, x  
    return 0, None
```

- **Steps:**

- Peeks at the heap's top entry (O(1)).
- If the entry is invalid (stock deleted or UUID outdated), pops it (O(log n)) and repeats.
- Returns the valid (-neg_vol, x) or (0, None) if empty.

- **Time Guarantee:** O(1) for valid top entry, O(log n) per pop for cleanup (amortized low due to infrequent updates).

Implementation Notes

- **BST Design:** The BST uses price as the primary key and id to break ties, ensuring unique nodes. The implementation avoids balancing (e.g., AVL or Red-Black Tree) for simplicity, assuming random insertions approximate balance (O(log n) average case).
- **Heap Cleanup:** The max heap uses lazy deletion (checking UUIDs in max_vol), reducing the need for immediate updates in increase_volume.
- **Hash Table Efficiency:** Python's dictionary ensures O(1) for lookup_by_id and other hash-based operations, critical for frequent lookups in testing.